

# Assignment2\_sub

September 25, 2018

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn import metrics
from datetime import datetime
from math import sin, cos, sqrt, atan2, radians
```

```
In [2]: # First lets read the data into table format so that we can draw some insight
# Next we will clean the data based on some observations.
data = pd.read_csv('train.csv', nrows=10000000)
```

```
In [3]: # Lets see what kind of data we have
data.describe()
```

```
Out[3]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	1.000000e+07	1.000000e+07	1.000000e+07	9.999931e+06	
mean	1.133854e+01	-7.250775e+01	3.991934e+01	-7.250897e+01	
std	9.799930e+00	1.299421e+01	9.322539e+00	1.287532e+01	
min	-1.077500e+02	-3.439245e+03	-3.492264e+03	-3.426601e+03	
25%	6.000000e+00	-7.399207e+01	4.073491e+01	-7.399139e+01	
50%	8.500000e+00	-7.398181e+01	4.075263e+01	-7.398016e+01	
75%	1.250000e+01	-7.396710e+01	4.076712e+01	-7.396367e+01	
max	1.273310e+03	3.457626e+03	3.344459e+03	3.457622e+03	


	dropoff_latitude	passenger_count
count	9.999931e+06	1.000000e+07
mean	3.991913e+01	1.684793e+00
std	9.237280e+00	1.323423e+00
min	-3.488080e+03	0.000000e+00
25%	4.073403e+01	1.000000e+00
50%	4.075316e+01	1.000000e+00
75%	4.076810e+01	2.000000e+00
max	3.351403e+03	2.080000e+02

```
In [4]: data[(data['passenger_count'] > 8) | (data['passenger_count'] <= 0)].shape
```

Out[4]: (35347, 8)


```
In [5]: # First lets check the variation of latitude and longitude by drawing the scatter plot t
# very relaxed boundary conditions
lat_border = (39, 42.5)
long_border = (-72.03, -75.75)
data.plot(kind='scatter', x='pickup_longitude', y='pickup_latitude',
          color='blue', s=.02, alpha=.6)
plt.title("Pickups")
plt.ylim(lat_border)
plt.xlim(long_border)
```

Out[5]: (-72.03, -75.75)




```
In [6]: lat_border = (39, 42.5)
long_border = (-72.03, -75.75)
data.plot(kind='scatter', x='dropoff_longitude', y='dropoff_latitude',
          color='green', s=.02, alpha=.6)
plt.title("Dropoffs")
plt.ylim(lat_border)
plt.xlim(long_border)
```

Out[6]: (-72.03, -75.75)




```
In [7]: #Plotting histograms as well to take a more closer look
        data[(data.pickup_latitude > 40) & (data.pickup_latitude < 42)].pickup_latitude.hist(bin
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1a10a306d8>
```




```
In [8]: data[(data.dropoff_latitude > 40) & (data.dropoff_latitude < 42)].dropoff_latitude.hist()
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x1a11401f60>
```




```
In [9]: data[(data.pickup_longitude > -74.5) & (data.pickup_longitude < -72)].pickup_longitude.h
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1a129c60f0>
```




```
In [10]: data[(data.dropoff_longitude > -74.5) & (data.dropoff_longitude < -72)].dropoff_longitude
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1a12915d30>
```



```
In [11]: data['passenger_count'].hist(bins=64, figsize=(16,8))
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1305db00>
```



```

In [12]: # This method will be used to do the initial phase opf the data cleaning,
# primarily outlier removals and some approximations.
def clean_data(df):
    # New york city has a central location cordinates of Latitude = 40.7128 and Longitu
    # Therefore we restrict the pickup and dropoff locations to avoid considering rides
    # are lying outside the nyc area.
    nyc_min_latitude = 40.45
    nyc_max_latitude = 40.97
    nyc_min_longitude = -74.28
    nyc_max_longitude = -73.64

    # Removing null entries from the data
    df = df.dropna(how='any', axis = 'rows')

    # Removing entries for which pickup/ dropoff locations do not lie inside the nyc ar
    df = df[(df['pickup_latitude'] >= nyc_min_latitude) & (df['pickup_latitude'] <= nyc
    df = df[(df['pickup_longitude'] >= nyc_min_longitude) & (df['pickup_longitude'] <=
    df = df[(df['dropoff_latitude'] >= nyc_min_latitude) & (df['dropoff_latitude'] <= n
    df = df[(df['dropoff_longitude'] >= nyc_min_longitude) & (df['dropoff_longitude'] <

    # Clean entries which have passenger count greater than 8
    df = df[(df['passenger_count'] > 0) & (df['passenger_count'] <= 8)]
    df = df[(df['fare_amount'] > 0) & (df['fare_amount'] < 100)]
    return df

In [13]: # Lets clean the data now to remove outliers and reduce our data to meaningful entries
data = clean_data(data)

In [14]: data.shape

Out[14]: (9742396, 8)

In [15]: stat = data.describe()

In [16]: stat[['pickup_latitude', 'dropoff_latitude', 'pickup_longitude', 'dropoff_longitude']]

Out[16]:
      pickup_latitude  dropoff_latitude  pickup_longitude  dropoff_longitude
count      9.742396e+06      9.742396e+06      9.742396e+06      9.742396e+06
mean        4.075084e+01      4.075122e+01     -7.397550e+01     -7.397458e+01
std         2.695317e-02      3.080515e-02      3.453672e-02      3.400300e-02
min         4.045191e+01      4.045191e+01     -7.427998e+01     -7.427996e+01
25%         4.073655e+01      4.073559e+01     -7.399229e+01     -7.399158e+01
50%         4.075334e+01      4.075385e+01     -7.398211e+01     -7.398062e+01
75%         4.076751e+01      4.076838e+01     -7.396836e+01     -7.396540e+01
max         4.096982e+01      4.096999e+01     -7.364037e+01     -7.364002e+01

In [17]: # There has been some reduction in the dataset. (Irrelevant data)
data.shape

Out[17]: (9742396, 8)

```



```
In [18]: # This code has been took from stack_overflow and gives the
# haversine distance between two points on the earth.
```

```
def get_euclidean_dist(loc_data):
    orig_lat , orig_lon, dest_lat, dest_lon = loc_data

    radius = 6371      # This is a constant whose value is equal to Earth's radius

    deltaLat = radians(dest_lat-orig_lat)
    deltaLon = radians(dest_lon-orig_lon)

    a = sin(deltaLat/2)**2 + cos(radians(orig_lat)) * cos(radians(dest_lat)) * sin(deltaLon/2)**2
    c = 2 * atan2(sqrt(a), sqrt(1-a))
    d = radius * c
    return d
```

```
In [19]: data.head()
```

```
Out[19]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21 UTC	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16 UTC	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00 UTC	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42 UTC	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00 UTC	

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	
4	-73.968095	40.768008	-73.956655	40.783762	

	passenger_count
0	1
1	1
2	2
3	1
4	1

```
In [20]: # Calculate the distance of each ride and make another entry in the dataset. This is an
# as fare amount is directly related to the distance travelled
columns = ['pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude']
data['distance'] = data[columns].apply(get_euclidean_dist, axis=1)
```

```
In [21]: data.head(5)
```

```
Out[21]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21 UTC	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16 UTC	

2	2011-08-18 00:35:00.000000049	5.7	2011-08-18 00:35:00 UTC
3	2012-04-21 04:30:42.00000001	7.7	2012-04-21 04:30:42 UTC
4	2010-03-09 07:51:00.0000000135	5.3	2010-03-09 07:51:00 UTC

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	
4	-73.968095	40.768008	-73.956655	40.783762	

	passenger_count	distance
0	1	1.030764
1	1	8.450134
2	2	1.389525
3	1	2.799270
4	1	1.999157

```
In [22]: data['pickup_datetime'] = pd.to_datetime(data['pickup_datetime'], format='%Y-%m-%d %H:%M:%S')
```

```
In [23]: pdt = data['pickup_datetime']
```

```
In [24]: dtIdx = pd.DatetimeIndex(data['pickup_datetime'])
```

```
In [25]: ## This function will give us the absolute time of the day in minutes.
## This is done to see if incorporating minutes in the time improves the correlation
def get_time_of_day(dateTime):
    dtIdx = pd.DatetimeIndex(dateTime)
    hours = dtIdx.hour
    minutes = dtIdx.minute
    absTimeofDay = (60*hours) + minutes
    return absTimeofDay
```

```
In [26]: data['time'] = get_time_of_day(data['pickup_datetime'])
```

```
In [27]: data.head()
```

```
Out[27]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.00000001	4.5	2009-06-15 17:26:21	
1	2010-01-05 16:52:16.00000002	16.9	2010-01-05 16:52:16	
2	2011-08-18 00:35:00.000000049	5.7	2011-08-18 00:35:00	
3	2012-04-21 04:30:42.00000001	7.7	2012-04-21 04:30:42	
4	2010-03-09 07:51:00.0000000135	5.3	2010-03-09 07:51:00	

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	

```

4          -73.968095          40.768008          -73.956655          40.783762

passenger_count  distance  time
0                1  1.030764  1046
1                1  8.450134  1012
2                2  1.389525    35
3                1  2.799270   270
4                1  1.999157   471

```

```


In [28]: ## Now we will see the relationship between the distance and the fare amount
## Since it is logical that a shorter ride will cost less than a longer one,
# based on this data we can actually get some insightful information
data.plot(kind='scatter',x='distance',y='fare_amount', s=0.2, alpha=0.4)

```

```

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x1a10a1c160>

```



```

In [29]: # Preprocessing data to fetch some useful information regarding the time of the ride/ t
data['year'] = pd.to_datetime(data['pickup_datetime']).dt.year
data['month'] = pd.to_datetime(data['pickup_datetime']).dt.month
data['day'] = pd.to_datetime(data['pickup_datetime']).dt.day
data['hour'] = pd.to_datetime(data['pickup_datetime']).dt.hour

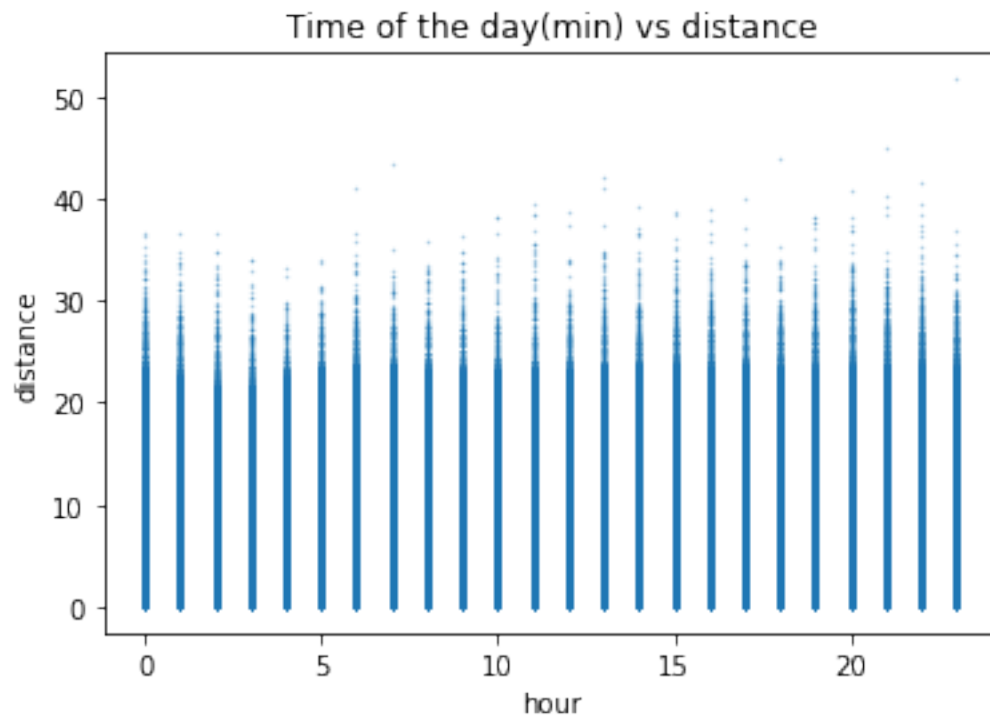
```

```

In [30]: #
data.plot(kind='scatter',x='hour',y='distance', s=0.2, alpha=0.4)
plt.title("Time of the day(min) vs distance")

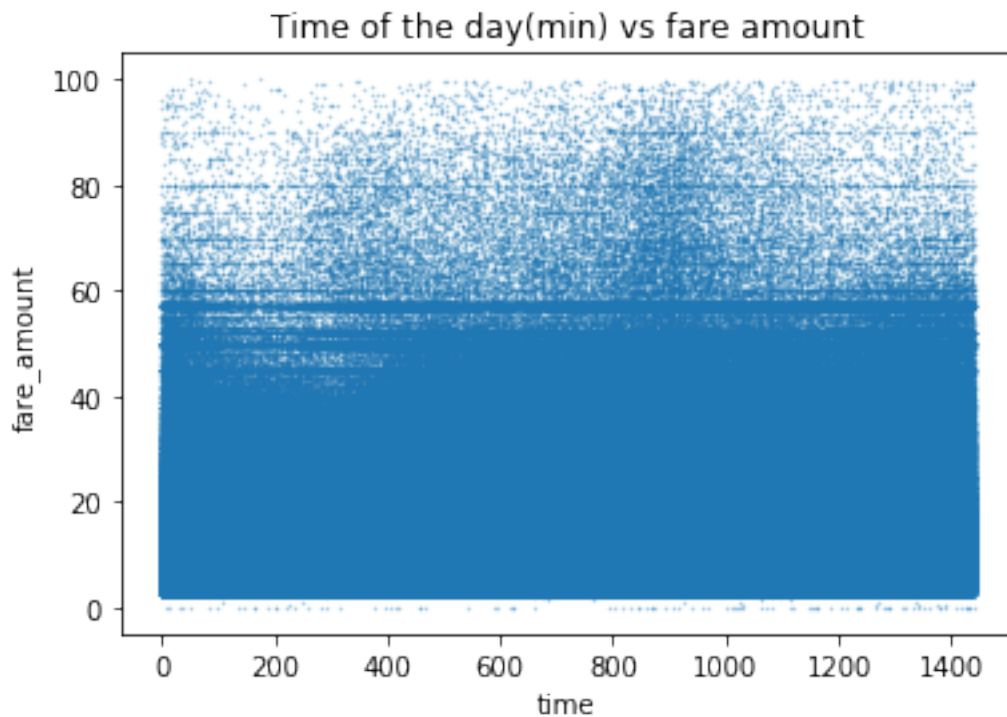
```

```
Out[30]: Text(0.5,1,'Time of the day(min) vs distance')
```



```
In [31]: data.plot(kind='scatter',x='hour',y='fare_amount', s=0.2, alpha=0.4)
plt.title("Time of the day(min) vs fare amount")
```

```
Out[31]: Text(0.5,1,'Time of the day(min) vs fare amount')
```



```
In [32]: data[['fare_amount', 'distance', 'time', 'hour']].corr()
```

```
Out[32]:
```

	fare_amount	distance	time	hour
fare_amount	1.000000	0.875650	-0.017688	-0.017387
distance	0.875650	1.000000	-0.029439	-0.029175
time	-0.017688	-0.029439	1.000000	0.999019
hour	-0.017387	-0.029175	0.999019	1.000000

```
In [33]: data['distance'].corr(data['fare_amount'])
```

```
Out[33]: 0.8756501057696913
```

```
In [34]: data.head()
```

```
Out[34]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00	

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	

2	-73.982738	40.761270	-73.991242	40.750562
3	-73.987130	40.733143	-73.991567	40.758092
4	-73.968095	40.768008	-73.956655	40.783762


	passenger_count	distance	time	year	month	day	hour
0	1	1.030764	1046	2009	6	15	17
1	1	8.450134	1012	2010	1	5	16
2	2	1.389525	35	2011	8	18	0
3	1	2.799270	270	2012	4	21	4
4	1	1.999157	471	2010	3	9	7

```
In [35]: data['rate'] = data['fare_amount']/data['distance']
```

```
In [36]: data = data[(data['rate'] > 0.5) & (data['rate'] < 10)]
```

```
In [37]: data.plot(kind='scatter',x='distance',y='fare_amount', s=0.2, alpha=0.4)
```

```
Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x1a12ab0ba8>
```




```
In [72]: data[['fare_amount','distance','time', 'hour']].corr()
```

```
Out[72]: 0.9261601234102027
```

```
In [39]: data.shape
```


```
Out[39]: (9308183, 15)
```

```
In [40]: # Lets see if there is some relationship b/w the day of the week and number of rides ta
# Usually more taxi's are booked during weekday.
data['day'].hist(bins=100, figsize=(16,8))
data['weekday'] = pd.to_datetime(data['pickup_datetime']).dt.weekday
```



```
In [41]: data['weekday'].hist(bins=32, figsize=(16,8))
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x1a13e43080>
```



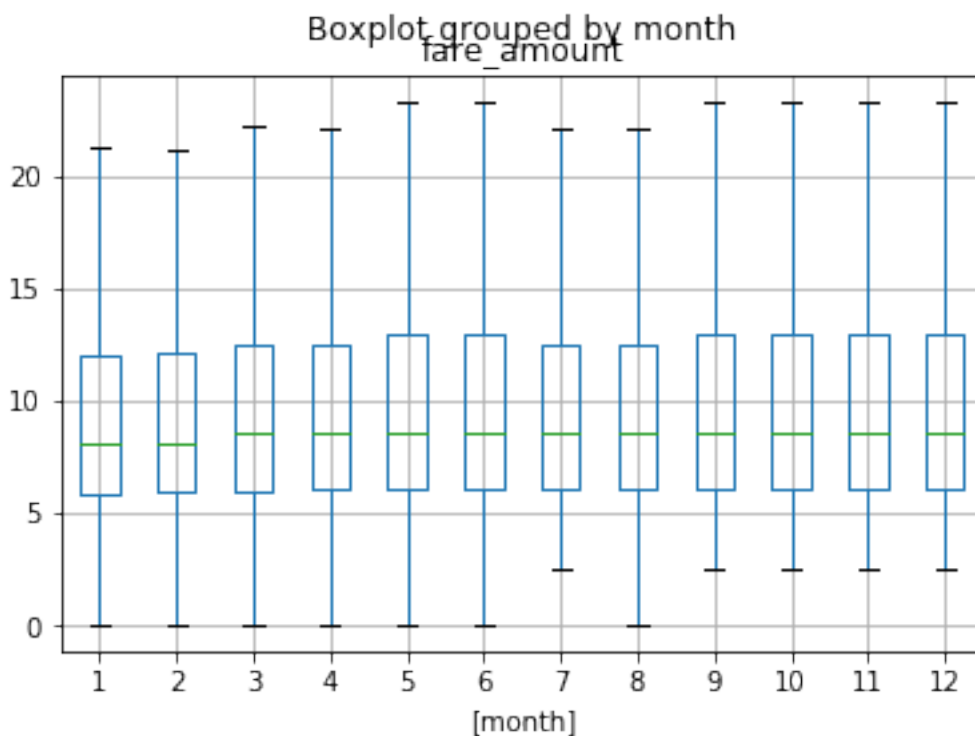
```
In [42]: # Plotting boxplot to get information regarding the distribution of data . Plotting for
# 1.Finding out if there exist a relation between the hour of taxi pickup and fare amou
```

```
# 2.Finding out if there exist a relation between the month of taxi pickup and fare amount
# 3.Finding out if there exist a relation between the year of taxi pickup and fare amount
```



```
data[['fare_amount', 'month']].boxplot(by='month', showfliers=False)
data[['fare_amount', 'day']].boxplot(by='day', showfliers=False)
data[['distance', 'hour']].boxplot(by='hour', showfliers=False)
data[['fare_amount', 'hour']].boxplot(by='hour', showfliers=False)
```


```
# From the below plot we can see that the fare is quite high in the morning hours, this
# airport rides as the same relation exist for the hour vs distance plot. Lets try to find
# actually corresponding to airport pickups and drops.
```

```
Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16b6b0f0>
```









```
In [43]: # Here we check if the given coordinates of the pickup location or dropoff location is m
# with the airports near the new york city. There are 3 airports near the city:
# 1. JFK
# 2. Laguardia
# 3. Newark
# Getting coordinates of these airports from the web
# JFK -> Latitude: 40.6413111, Longitude: -73.7781391
# LaG -> Latitude: 40.77725, Longitude: -73.872611
# Newark -> Latitude: 40.6925, Longitude: -74.168611
```

```
def get_is_airport_ride(column):
    pickup_lat, pickup_lon, dropoff_lat, dropoff_lon = column
    jfk_airport = (40.6413, -73.778)
    lag_airport = (40.777, -73.872)
    newrk_airport = (40.692, -74.168)

    is_drop_at_jkf = (pickup_lat, pickup_lon, jfk_airport[0], jfk_airport[1] )
    is_drop_at_lag = (pickup_lat, pickup_lon, lag_airport[0], lag_airport[1] )
    is_drop_at_newrk = (pickup_lat, pickup_lon, newrk_airport[0], newrk_airport[1] )

    is_pickup_from_jkf = (jfk_airport[0], jfk_airport[1], dropoff_lat, dropoff_lon )
    is_pickup_from_lag = (lag_airport[0], lag_airport[1], dropoff_lat, dropoff_lon )
```

```

is_pickup_from_newrk = (newrk_airport[0], newrk_airport[1], dropoff_lat, dropoff_lo

if(get_euclidean_dist(is_pickup_from_jkf) < 1 or get_euclidean_dist(is_pickup_from_
    get_euclidean_dist(is_pickup_from_newrk) < 1 ):
    return 1
if(get_euclidean_dist(is_drop_at_jkf) < 1 or get_euclidean_dist(is_drop_at_lag) < 1
    get_euclidean_dist(is_drop_at_newrk) < 1 ):
    return 0

return 0

```

```

In [44]: # Adding a new feature, which will give us information about airport rides
# Based on this information, I plan to do further analysis on the pattern of long dista
columns = ['pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude']
data['is_airport_ride'] = data[columns].apply(get_is_airport_ride, axis=1)

```

```


In [45]: data['is_airport_ride'].hist(bins=32, figsize=(16,8))

```

```

Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x1a10d06da0>

```



```

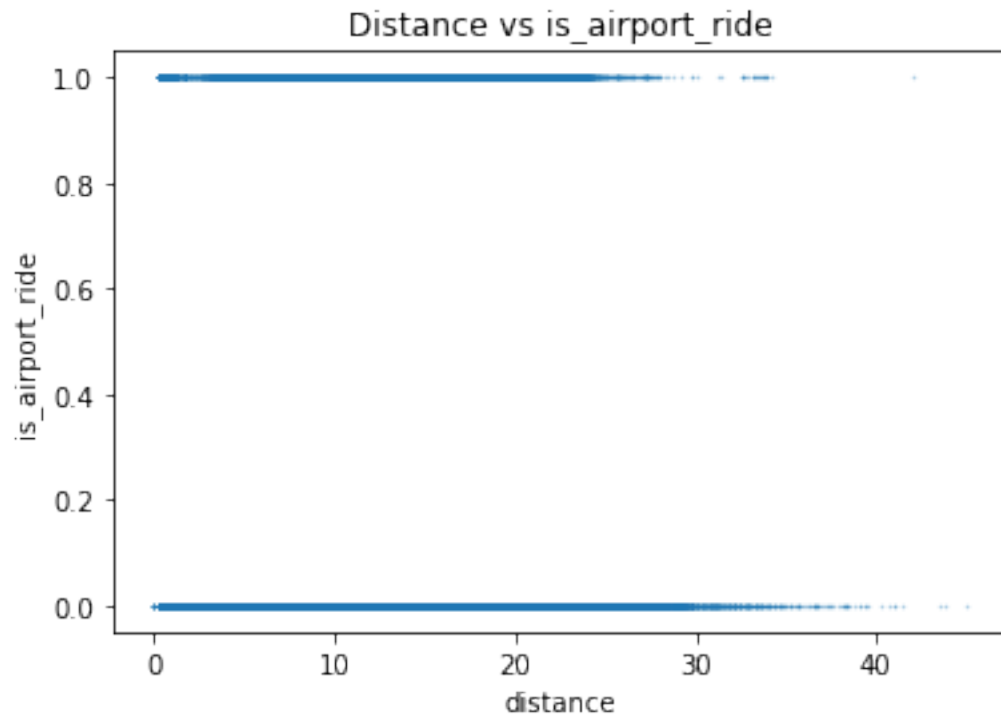
In [46]: data.plot(kind='scatter',x='distance',y='is_airport_ride', s=0.2, alpha=0.4)
plt.title("Distance vs is_airport_ride")

```

```


Out[46]: Text(0.5,1,'Distance vs is_airport_ride')

```



```
In [47]: # Lets find out if there is something interesting we can do with the data we collected
# and the hour of the day. This plot will help us analyze the airport taxi traffic thro
fig, ax = plt.subplots(figsize=(15,7))
data_airport = data[data['is_airport_ride'] > 0]
data_airport.groupby(['hour', 'is_airport_ride']).count()['key'].unstack().plot(ax=ax)
```

```
Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16b71b38>
```



```

In [48]: features_to_keep = ['pickup_latitude', 'pickup_longitude', 'dropoff_longitude',
                             'dropoff_latitude', 'passenger_count', 'distance', 'hour', 'year', 'is_airport_ride']

In [49]: train_data = data[features_to_keep]

In [50]: train_data.head()

Out[50]:
  pickup_latitude  pickup_longitude  dropoff_longitude  dropoff_latitude  \
0      40.721319      -73.844311      -73.841610      40.712278
1      40.711303      -74.016048      -73.979268      40.782004
2      40.761270      -73.982738      -73.991242      40.750562
3      40.733143      -73.987130      -73.991567      40.758092
4      40.768008      -73.968095      -73.956655      40.783762

  passenger_count  distance  hour  year  is_airport_ride
0                1  1.030764    17  2009                0
1                1  8.450134    16  2010                0
2                2  1.389525     0  2011                0
3                1  2.799270     4  2012                0
4                1  1.999157     7  2010                0

In [51]: output_data = data['fare_amount']

In [52]: from sklearn.linear_model import LinearRegression
          from sklearn.metrics import mean_squared_error
          from sklearn.model_selection import train_test_split

In [53]: X_train, X_test, y_train, y_test = train_test_split(train_data, output_data, test_size=0.2)

In [54]: linear_reg = LinearRegression()

In [55]: linear_reg.fit(X_train, y_train)

Out[55]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [56]: print(linear_reg.coef_)

[ 6.21710727e+00  1.64002304e+01 -1.49944810e+01 -1.05592095e+01
  3.46015120e-02  2.28005761e+00  1.36026274e-02  5.09205560e-01
  6.70096628e+00]

In [57]: y_pred = linear_reg.predict(X_test)

In [58]: rmse = np.sqrt(mean_squared_error(y_test, y_pred))
          print("root mean squared error: {}".format(rmse))

```

root mean Squared error: 3.202870444918603

```
In [59]: test_data = pd.read_csv('test.csv')
```

```
In [60]: test_data.head()
```

```
Out [60]:
```

		key	pickup_datetime	pickup_longitude \
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24 UTC	-73.973320	
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24 UTC	-73.986862	
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44 UTC	-73.982524	
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12 UTC	-73.981160	
4	2012-12-01 21:12:12.0000003	2012-12-01 21:12:12 UTC	-73.966046	

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	40.763805	-73.981430	40.743835	1
1	40.719383	-73.998886	40.739201	1
2	40.751260	-73.979654	40.746139	1
3	40.767807	-73.990448	40.751635	1
4	40.789775	-73.988565	40.744427	1

```
In [61]: test_data.count()
```

```
Out [61]: key          9914
pickup_datetime      9914
pickup_longitude     9914
pickup_latitude      9914
dropoff_longitude    9914
dropoff_latitude     9914
passenger_count      9914
dtype: int64
```

```
In [62]: # just like training data , we add the 'distance' field to the test data as part of pre
columns = ['pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude']
test_data['distance'] = test_data[columns].apply(get_euclidean_dist, axis=1)
test_data['is_airport_ride'] = test_data[columns].apply(get_is_airport_ride, axis=1)
```

```
In [63]: # Making test data same as training data in terms of representation so that the model
# doesn't complain.
test_data['hour'] = test_data.pickup_datetime.apply(lambda t: pd.to_datetime(t).hour)
test_data['year'] = test_data.pickup_datetime.apply(lambda t: pd.to_datetime(t).year)
# Keeping this as it is required to be output in the submission.csv
test_data_with_key = test_data[['key'] + features_to_keep]
test_data = test_data[features_to_keep]
```

```
In [64]: test_data.head()
```

```
Out [64]:
```

	pickup_latitude	pickup_longitude	dropoff_longitude	dropoff_latitude \
0	40.763805	-73.973320	-73.981430	40.743835

1	40.719383	-73.986862	-73.998886	40.739201
2	40.751260	-73.982524	-73.979654	40.746139
3	40.767807	-73.981160	-73.990448	40.751635
4	40.789775	-73.966046	-73.988565	40.744427

	passenger_count	distance	hour	year	is_airport_ride
0	1	2.323260	13	2015	0
1	1	2.425353	13	2015	0
2	1	0.618628	11	2011	0
3	1	1.961033	21	2012	0
4	1	5.387301	21	2012	0

```
In [65]: # Getting the prediction results from the linear regressor model and output it to the s
linear_reg.fit(train_data, output_data)
test_predictions = linear_reg.predict(test_data)
```

```
In [66]: len(test_predictions)
```

```
Out[66]: 9914
```

```
In [67]: submission = pd.DataFrame(
    {'key': test_data_with_key.key, 'fare_amount': test_predictions},
    columns = ['key', 'fare_amount'])
```

```
In [68]: submission.to_csv('submission.csv', index = False)
```

```
In [69]: # Got a score of 5.35 with k-fold Linear regression
# Now trying random Forest regressor to check if there is any improvement
from sklearn.ensemble import RandomForestRegressor
rfgModel = RandomForestRegressor()
# Trying cross validation first to check if the model is givibng good results. Root mea
# a good approximation of the performance of a prediction model
print("Random Forest Generator Parameters: ")
print(rfgModel.get_params() )
rfgModel.fit(X_train, y_train)
rfgModel_pred = rfgModel.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test , rfgModel_pred))
print("root mean Squared error: {}".format(rmse))
rfgModel.fit(train_data, output_data)
# Now running the model on actual data test data
rfgModel_pred = rfgModel.predict(test_data)
```

Random Forest Generator Parameters:

```
{'bootstrap': True, 'criterion': 'mse', 'max_depth': None, 'max_features': 'auto', 'max_leaf_nod
root mean Squared error: 2.4442165252751806
```

```
In [70]: submission = pd.DataFrame(
    {'key': test_data_with_key.key, 'fare_amount': rfgModel_pred},
    columns = ['key', 'fare_amount'])
```

```
In [71]: submission.to_csv('submission.csv', index = False)
```